

# Cache Memory

---

- What is cache memory?
- Cache Organisation
- Cache Implementation
- Cache Performance

# Cache Memory

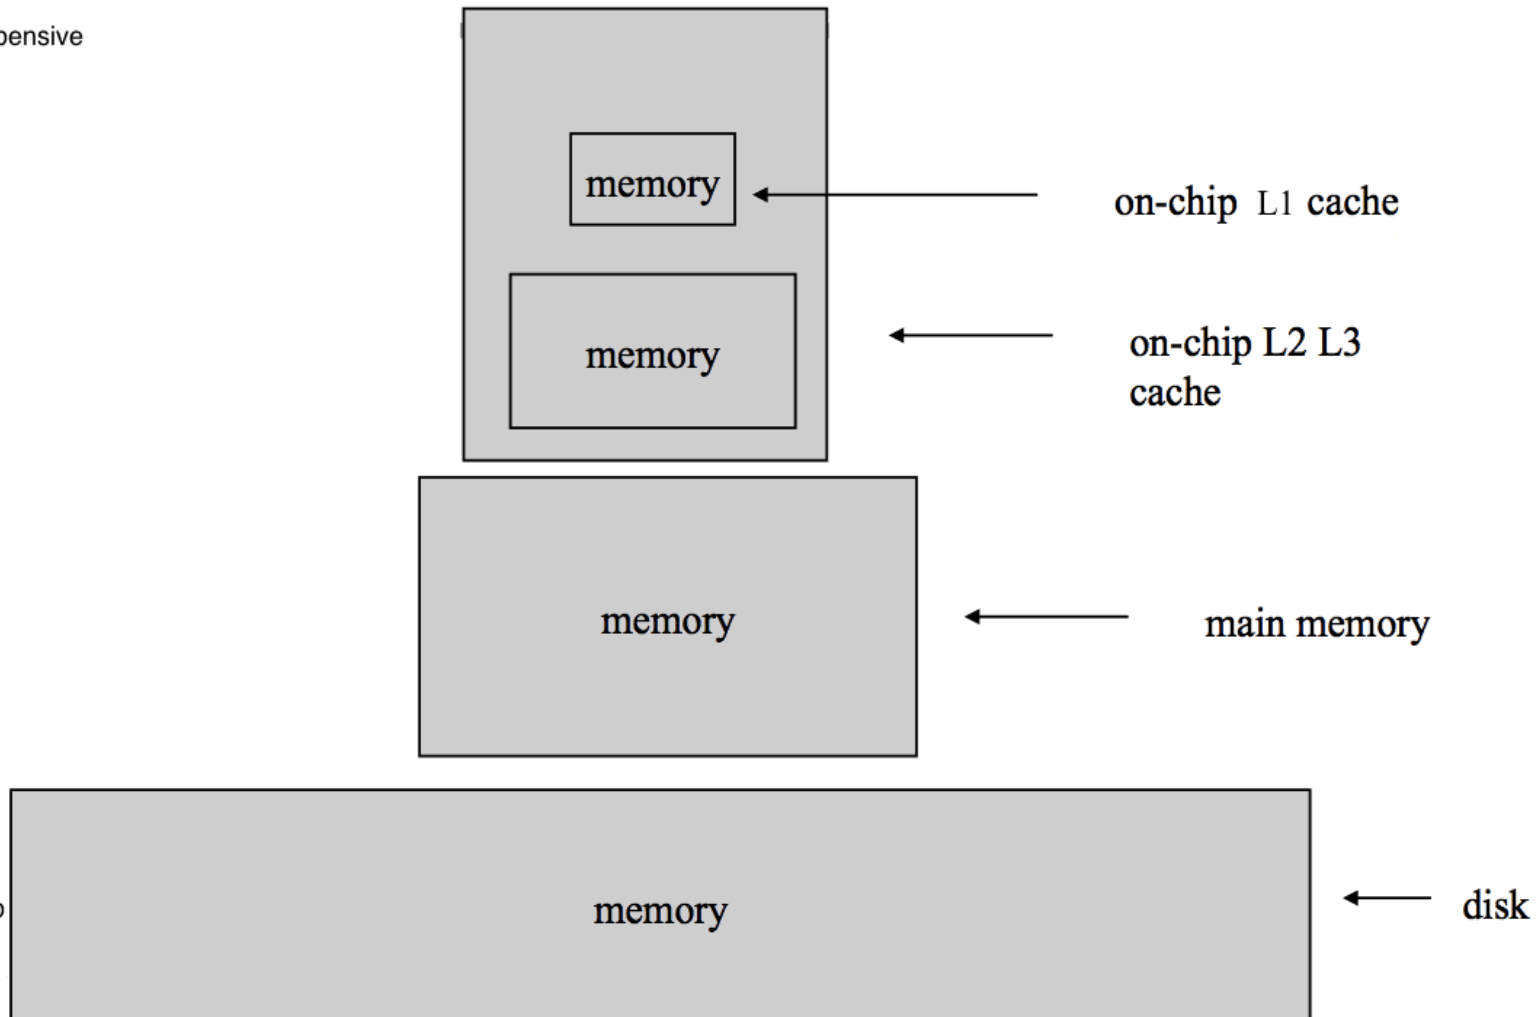
---

- Early processors had an on-chip instruction queue, a modest amount of DRAM and a hard disk
- As processor speeds increased, memory access times did not keep pace and more creative solutions were required for memory interfaces.
- A multi-level memory architecture known as ‘cache’ provided a solution.

Small / Fast / Expensive



Big / Slow / Cheap

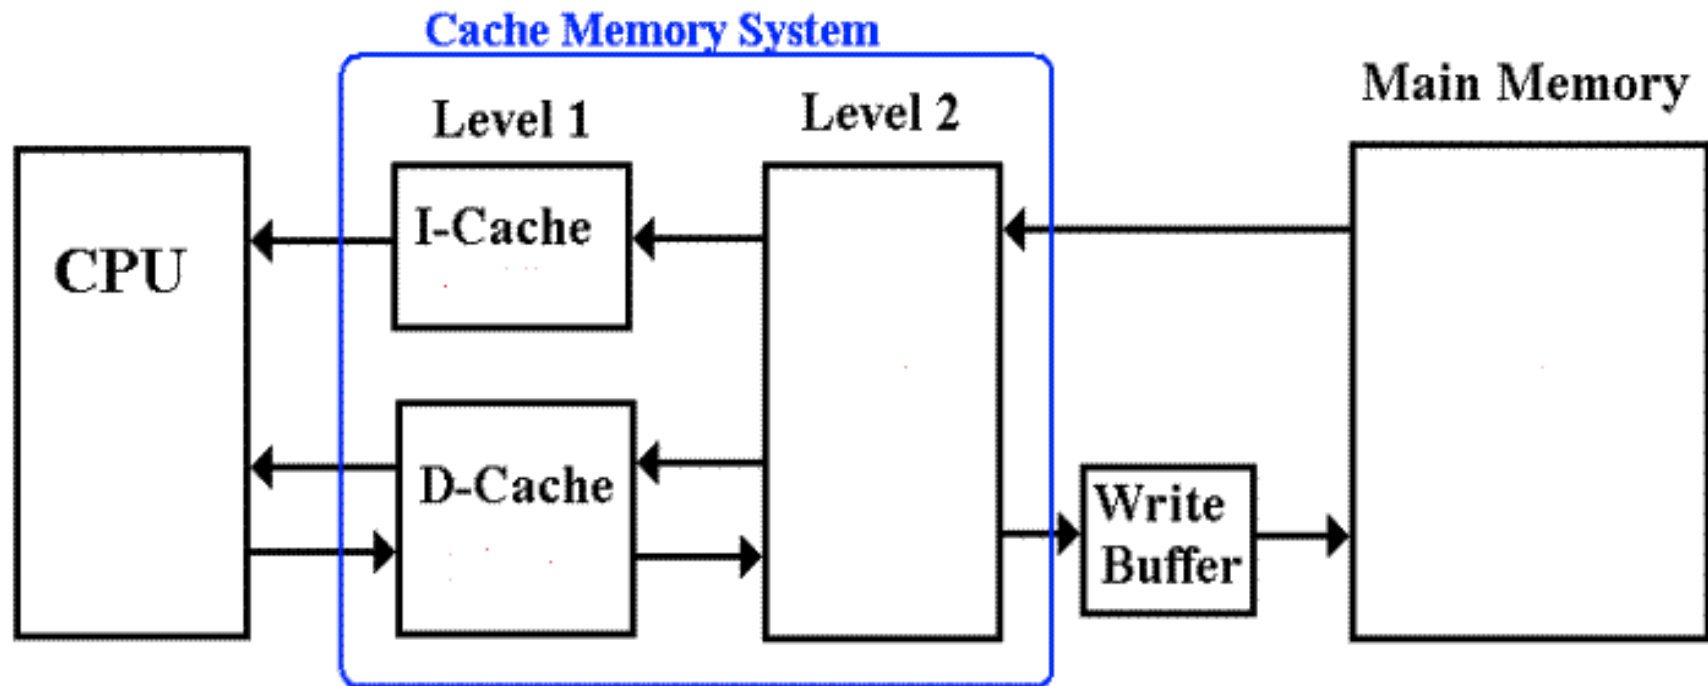


# Cache Memory

---


- Cache memory is a temporary store for data and instructions; it is faster to access than main memory
- L1 cache is a small memory on the processor chip which runs at processor speed
- L2 cache is much larger than L1 but slower. It is accessed if the required data is not found in the L1 cache; aim for L2 to have the data about 90% of the time
- Multi-core chips can have a L3 cache shared by all of the cores

# Cache Memory



# Cache Memory

---



# Cache Definitions

---

- *cache hit* – cpu requested data is found in the next level cache
- *cache miss* – requested data is not in the cache and must be fetched from upper levels
- *hit time* -- time to access the cache (time to determine hit/miss + memory access time)
- *miss penalty* -- time to move data from upper levels to cpu
- *hit ratio* -- percentage of time the data is found in the cache
- *miss ratio* --  $(1 - \text{hit ratio})$

# Types of Organisation

---

- Direct Mapped
  - Fully Associative
  - Set Associative
  - + few others
- 
- Caches are arranged in blocks (or lines) just like lower levels of the hierarchy



# Direct Mapped

---

- Simplest and cheapest
- Lowest performance
- Some significant drawbacks

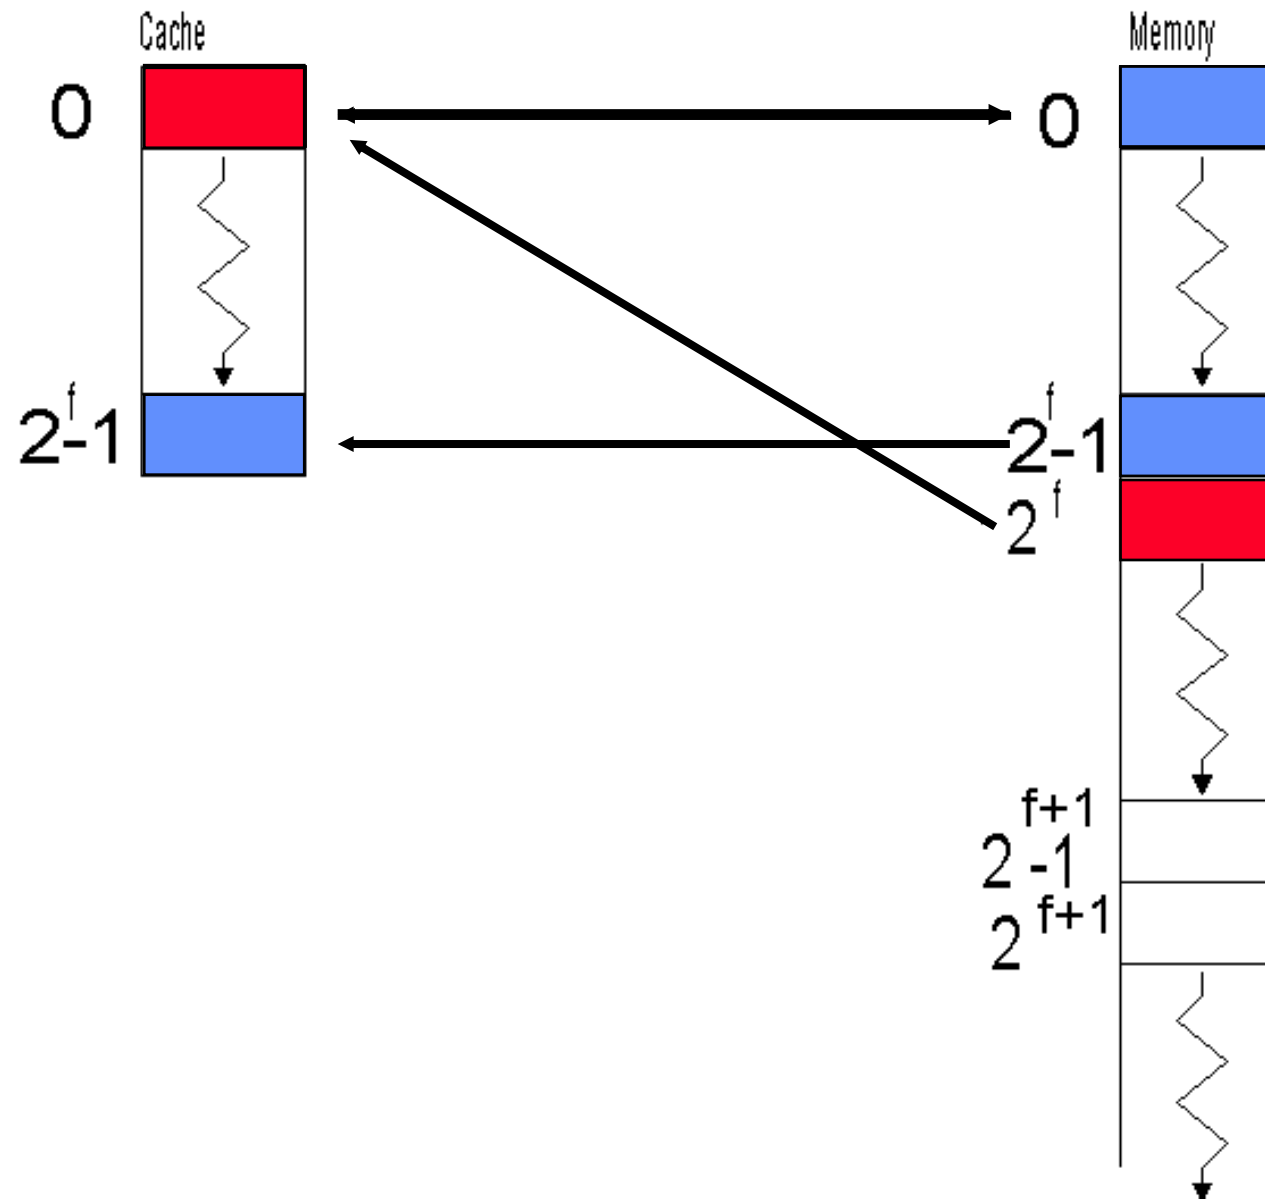
# How does it work

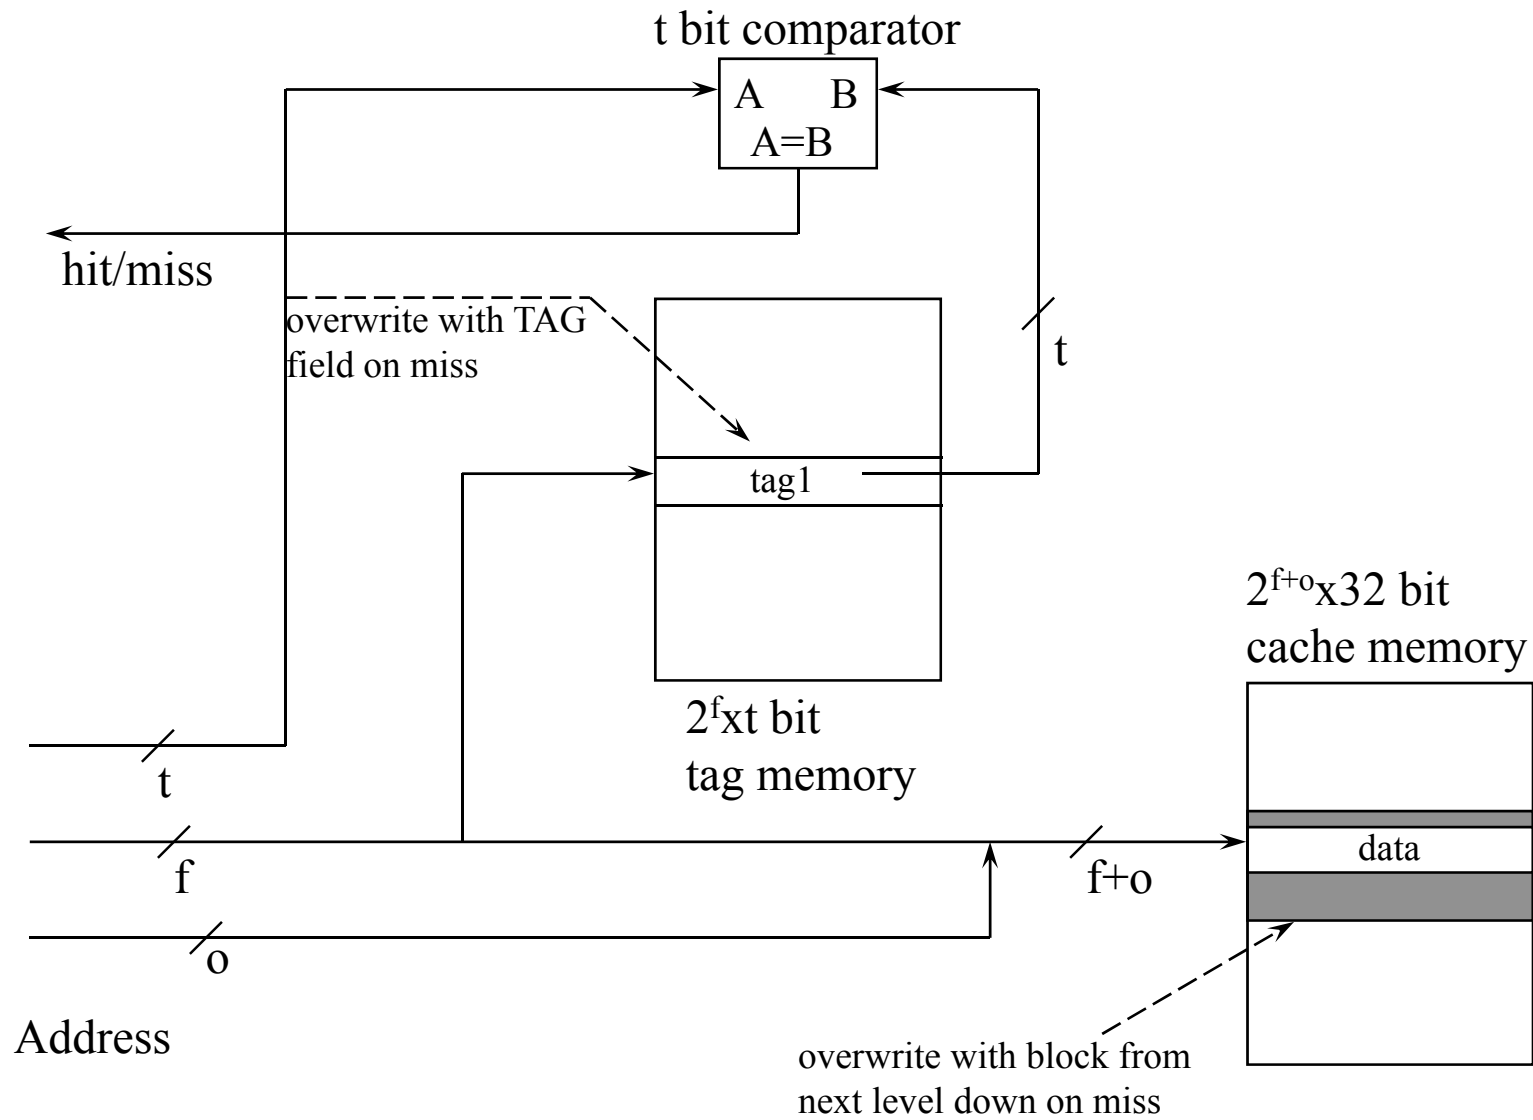
---

- Address to Mapper organised as follows

TAG	FRAME	OFFSET
-----	-------	--------

- OFFSET is offset to the cache block and passes through unchanged
- FRAME is address of block in cache and passes through unchanged
- TAG is used to identify whether the requested block is in the cache





# Advantages/Disadvantages

---

- Low hardware overhead
- Fast on hit
  - read doesn't add any time
  - On write overall access time is tag memory time + cache memory time
- Block from next level down can only be mapped into a *specific* block

# Fully Associative

---

- Block from next level down can be mapped into any block in the cache
  - Similar to Primary/Secondary memory interface

# Advantages/Disadvantages

---

- Very Flexible
- Does not suffer from problem due to blocks which cannot co-exist in the cache
- Needs full TLB to find block in cache
- Access time always incurs the overhead of accessing the TLB
- Big hardware overhead
- Needs memory management (RANDOM ?)

# Set Associative

---

- Combines advantages of both Direct Mapped and Fully Associative
  - Hardware is relatively simple
  - A Practical cache can be constructed which rivals the Fully Associative Cache in performance terms




# How does it Work

---


- Address is as follows

TAG	SET	OFFSET
-----	-----	--------

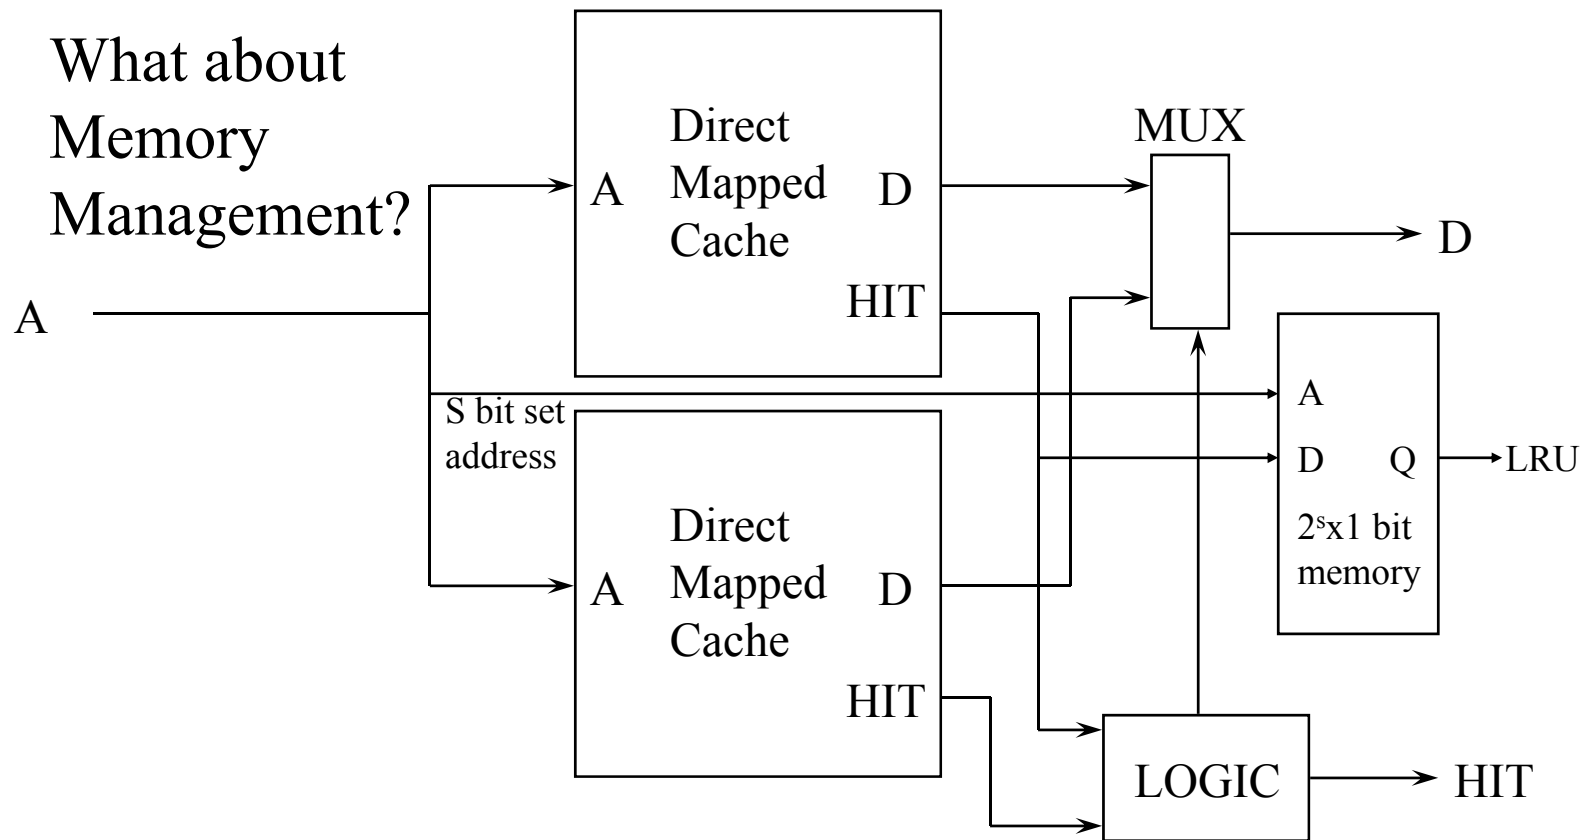
- Cache is organised into sets
- Each set contains a number of individual blocks or frames - say  $n$
- Cache is Direct Mapped with respect to sets (using SET and TAG fields)
- Cache is Fully Associative within a set



## n-way Set Associative Cache



# How is it Implemented?



## 2-way Set Associative Cache

# Complexity w.r.t. $n$

---

- As  $n$  becomes bigger, the cache looks more like a Fully Associative Cache
- 2-way set associative is easy just 2 Direct Mapped Caches with a bit of extra hardware to select between them
- What about memory management?

# • Can you explain this?

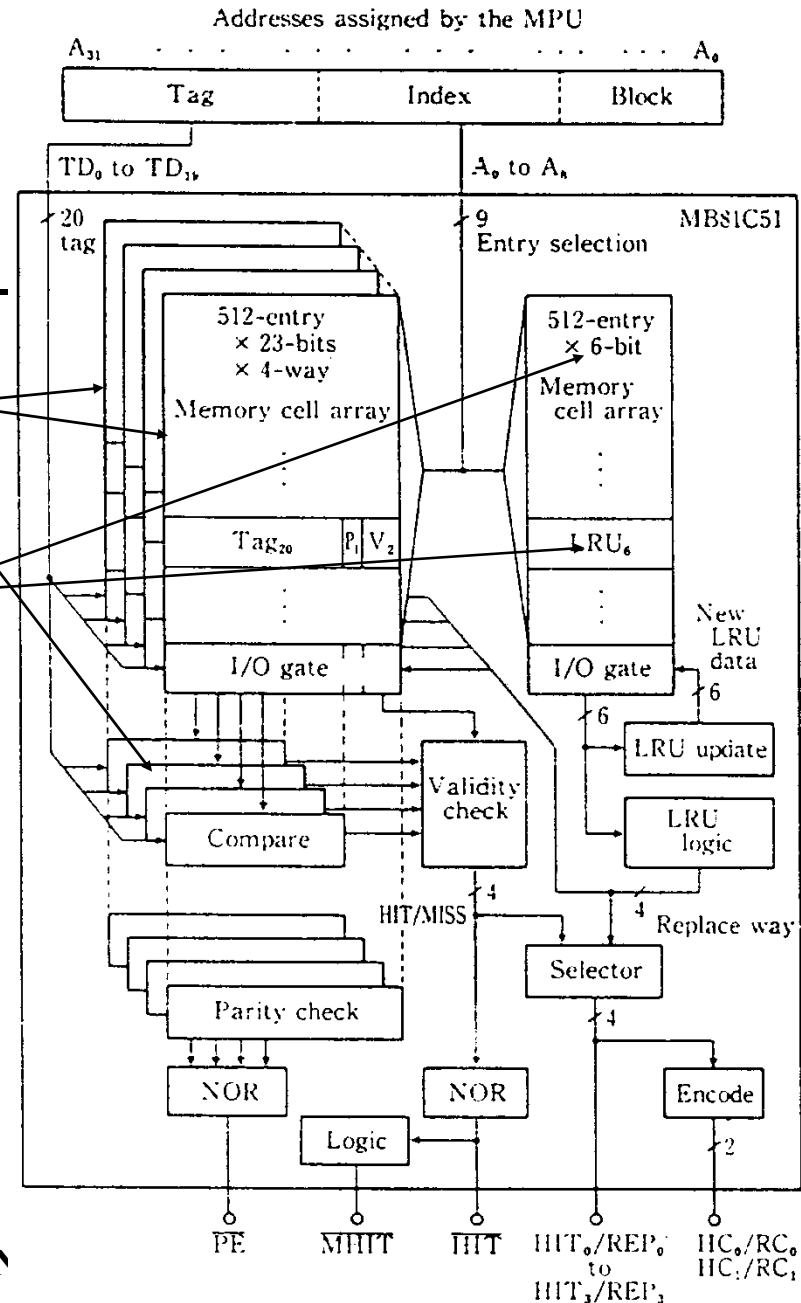
4-way SA (4 sets of LUTs)

Only one (or none) will HIT

LRU on each set

Why 6 bits / set

- 4 in each set
- 2 bit ID
- A Queue of 3 IDs is all that is necessary  
= 6 bits / set



# Other Issues


---

- Can separate caches for program and data
  - Program cache is never written
  - Caches can be optimised for underlying behaviour of data
  - Disadvantage - decision about balance between program and data is predetermined
- Got to be able to invalidate cache contents
- Some memory must be *uncacheable*

# Writing to a Cache

---

- Two options on writing to data cache:
  - Write back - When a block is discarded write it back to the next level down
  - Write through - When a datum is written write it to the next level down as well (but don't wait for the write to complete)



# Write-Through

---

- Means that blocks do not need to be written back but:
  - Can cause the processor to wait if a write is required before the previous write is finished
  - Two further options with write-through when there is a miss on a write
    - bring the block into the cache - With Allocate
    - don't bring the block - With No Allocate




# Cache Performance

---

- There are analytical models but, generally, performance must be derived empirically for representative work-loads
- Performance depends on work-load, size, type, line-size, etc.

# Case Study (L1 Cache)



C-L Su, A M Despain, Cache Design Trade-Offs for Power and Performance Optimisation: A Case Study. Proc. ISLPED, pp63-68, 1995

# Execution Trace Cache

---

- Intended for instructions
- They do not store an image of the memory
- They store the sequence of instructions that have been executed
- This implies a greater efficiency
- Instructions can be partly decoded before being put in the TEC

# Example

- Code snippet

```

L1  • move x,y
    • cmp y,#0
    • je L2
      inst 1
      inst 2
      inst 3
L2  • cmp y,#9
    • add x,#1
    • jlt L3
      inst 4
      inst 5
L3  • in x,23
    • jmp L1
    
```

## Normal cache contents

• move x,y	• cmp y,#0	• je L2	inst 1
inst 2	inst 3	• cmp y,#9	• add x,#1
• jlt L3	inst 4	• in x,23	• jmp L1

## Execution trace cache contents

• move x,y	• cmp y,#0	• je L2	• cmp y,#9
• add x,#1	• jlt L3	• in x,23	• jmp L1